

<https://kubernetes.io/docs/tutorials/kubernetes-basics/>

+Learning Kubernetes Basics

<https://kubernetes.io/docs/tutorials/kubernetes-basics/>

Notes to share with anyone interesting in taking above tutorial – see License at Kubernetes.io. Notes provided as is as courtesy for your review from <https://longislandnewtechnology.blogspot.com>

Companion notes to <https://kubernetes.io/docs/tutorials/kubernetes-basics>

Notes intended to be shared with anyone interesting in taking above tutorial -> see Tutorial License at Kubernetes.io.

Annotated Notes with screenshots provided "AS IS" courtesy of <https://longislandnewtechnology.blogspot.com>

PDF document is "annotated guide" from K8SSG (tm - longislandnewtechnology) aka Kubernetes Study Group of Long Island, New York, USA

Any screenshots used elsewhere require link back to blog address <https://longislandnewtechnology.blogspot.com>

+Learning Kubernetes Basics

<https://kubernetes.io/docs/tutorials/kubernetes-basics/>

Using Minikube to create a cluster

<https://kubernetes.io/docs/tutorials/kubernetes-basics/create-cluster/cluster-intro/>

CREATE A CLUSTER

Once minikube is installed as a local VM, check to see if minikube is installed properly: type *minikube version*

Kubernetes Bootcamp Terminal

```
$  
$ minikube version  
minikube version: v1.6.2  
commit: 54f28ac5d3a815d1196cd5d57d707439ee4bb392
```

Start the cluster by typing "minikube start" there will be a delay here ... when complete it will look like:

```
$ minikube start  
* minikube v1.6.2 on Ubuntu 18.04  
* Selecting 'none' driver from user configuration (alternates: [])  
* Running on localhost (CPUs=2, Memory=2461MB, Disk=47990MB) ...  
* OS release is Ubuntu 18.04.3 LTS  
* Preparing Kubernetes v1.17.0 on Docker '18.09.7' ...  
  - kubelet.resolv-conf=/run/systemd/resolve/resolv.conf  
  
* Pulling images ...  
* Launching Kubernetes ...  
* Configuring local host environment ...  
* Waiting for cluster to come online ...  
* Done! kubectl is now configured to use "minikube"  
$  
$
```

You can run some ad hoc kubectl commands if you know some...

```
$ kubectl get pod  
No resources found in default namespace.  
$ kubectl get service  


| NAME       | TYPE      | CLUSTER-IP | EXTERNAL-IP | PORT(S) | AGE |
|------------|-----------|------------|-------------|---------|-----|
| kubernetes | ClusterIP | 10.96.0.1  | <none>      | 443/TCP | 88s |

  
$
```

Module 2 – check your cluster version by typing **kubectl version**

```
$ kubectl version
Client Version: version.Info{Major:"1", Minor:"17", GitVersion:"v1.17.0", GitCommitt:"70132b0f130acc0bed193d9ba59dd186f0e634cf", GitTreeState:"clean", BuildDate:"2019-12-07T21:20:10Z", GoVersion:"go1.13.4", Compiler:"gc", Platform:"linux/amd64"}
Server Version: version.Info{Major:"1", Minor:"17", GitVersion:"v1.17.0", GitCommitt:"70132b0f130acc0bed193d9ba59dd186f0e634cf", GitTreeState:"clean", BuildDate:"2019-12-07T21:12:17Z", GoVersion:"go1.13.4", Compiler:"gc", Platform:"linux/amd64"}
$
```

The client version is the **kubectl version (v1.17.0)** while the server version is the **kubernetes version (v1.17.0)**

LEARN MORE ABOUT YOUR CLUSTER

kubectl cluster-info

```
$ kubectl cluster-info
Kubernetes master is running at https://172.17.0.10:8443
KubeDNS is running at https://172.17.0.10:8443/api/v1/namespaces/kube-system/services/kube-dns:dns/proxy

To further debug and diagnose cluster problems, use 'kubectl cluster-info dump'.
$
```

Powered by  Katacoda

Notes: Cluster has a single node. You can use CLI or Dashboard (the latter allows you to view your applications in a UI)

kubectl get nodes

```
$ kubectl get nodes
NAME          STATUS    ROLES    AGE   VERSION
minikube      Ready     master   15m   v1.17.0
```

A status of ready means “node is ready to accept applications for deployment”

Using kubectl to Create a Deployment

<https://kubernetes.io/docs/tutorials/kubernetes-basics/deploy-app/deploy-intro/>

<https://kubernetes.io/docs/tutorials/hello-minikube/>

```
$ sleep 1; launch.sh
```

```
# kubectl get nodes --help
```

```
$ kubectl get nodes --help
Display one or many resources
```

```
Prints a table of the most important information about the specified resources.
You can filter the list using a label selector and the --selector flag. If the
desired resource type is namespaced you will only see results in your current
namespace unless you pass --all-namespaces.
```

```
Uninitialized objects are not shown unless --include-uninitialized is passed.
```

```
By specifying the output as 'template' and providing a Go template as the value
of the --template flag, you can filter the attributes of the fetched resources.
```

```
Use "kubectl api-resources" for a complete list of supported resources.
```

```
Examples:
```

```
# List all pods in ps output format.
```

```
kubectl get pods
```

```
# List all pods in ps output format with more information (such as node name).
```

```
kubectl get pods -o wide
```

```
# List a single replication controller with specified NAME in ps output format.
```

```
kubectl get replicationcontroller web
```

```
# List deployments in JSON output format, in the "v1" version of the "apps" API group:
```

```
kubectl get deployments.v1.apps -o json
```

```
# List a single pod in JSON output format.
```

```
kubectl get -o json pod web-pod-13je7
```

```
# List a pod identified by type and name specified in "pod.yaml" in JSON output format.
```

```
kubectl get -f pod.yaml -o json
```

```
# List resources from a directory with kustomization.yaml - e.g. dir/kustomization.yaml.
```

```
kubectl get -k dir/
```

```
# Return only the phase value of the specified pod.
```

```
kubectl get -o template pod/web-pod-13je7 --template={{.status.phase}}
```

```
# List resource information in custom columns.
```

```
kubectl get pod test-pod -o custom-columns=CONTAINER:.spec.containers[0].name,IMAGE:.spec.c
ontainers[0].image
```

```
# List all replication controllers and services together in ps output format.
kubectl get rc,services
```

```
# List one or more resources by their type and names.
kubectl get rc/web service/frontend pods/web-pod-13je7
```

Options:

-A, --all-namespaces=false: If present, list the requested object(s) across all namespaces. Namespace in current context is ignored even if specified with --namespace.

--allow-missing-template-keys=true: If true, ignore any errors in templates when a field or map key is missing in the template. Only applies to goyaml and jsonpath output formats.

--chunk-size=500: Return large lists in chunks rather than all at once. Pass 0 to disable. This flag is beta and may change in the future.

--field-selector='': Selector (field query) to filter on, supports '=', '==', and '!='. (e.g. --field-selector key1=value1,key2=value2). The server only supports a limited number of field queries per type.

-f, --filename=[]: Filename, directory, or URL to files identifying the resource to get from a server.

--ignore-not-found=false: If the requested object does not exist the command will return exit code 0.

-k, --kustomize='': Process the kustomization directory. This flag can't be used together with -f or -R.

-L, --label-columns=[]: Accepts a comma separated list of labels that are going to be presented as columns. Names are case-sensitive. You can also use multiple flag options like -L label1 -L label2...

--no-headers=false: When using the default or custom-column output format, don't print headers (default print headers).

-o, --output='': Output format. One of: json|yaml|wide|name|custom-columns=...|custom-columns-file=...|go-template=...|go-template-file=...|jsonpath=...|jsonpath-file=... See custom columns [http://kubernetes.io/docs/user-guide/kubectl-overview/#custom-columns], goyaml template [http://golang.org/pkg/text/template/#pkg-overview] and jsonpath template [http://kubernetes.io/docs/user-guide/jsonpath].

--raw='': Raw URI to request from the server. Uses the transport specified by the kubeconfig file.

-R, --recursive=false: Process the directory used in -f, --filename recursively. Useful when you want to manage related manifests organized within the same directory.

-l, --selector='': Selector (label query) to filter on, supports '=', '==', and '!='. (e.g. -l key1=value1,key2=value2)

--server-print=true: If true, have the server return the appropriate table output. Supports extension APIs and CRDs.

--show-kind=false: If present, list the resource type for the requested object(s).

--show-labels=false: When printing, show all labels as the last column (default hide labels column)

--sort-by='': If non-empty, sort list types using this field specification. The field specification is expressed as a JSONPath expression (e.g. '{.metadata.name}'). The field in the API resource specified by this JSONPath expression must be an integer or a string.

--template='': Template string or path to template file to use when -o=go-template, -o=go-template-file. The template format is goyaml templates [http://golang.org/pkg/text/template/#pkg-overview].

-w, --watch=false: After listing/getting the requested object, watch for changes. Uninitialized objects are excluded if no object name is provided.

--watch-only=false: Watch for changes to the requested object(s), without listing/getting first.

Usage:

```
kubectl get [(-o|--output)=]json|yaml|wide|custom-columns=...|custom-columns-file=...|go-template=...|go-template-file=...|jsonpath=...|jsonpath-file=... (TYPE[.VERSION][.GROUP] [NAME | -l label] | TYPE[.VERSION][.GROUP]/NAME ...) [flags] [options]
```

Use "kubectl options" for a list of global command-line options (applies to all commands).

```
$
```

\$ kubectl version

```
$ kubectl version
Client Version: version.Info{Major:"1", Minor:"15", GitVersion:"v1.15.2", GitCommit:"f6278300be
bbb750328ac16ee6dd3aa7d3549568", GitTreeState:"clean", BuildDate:"2019-08-05T09:23:26Z", GoVers
ion:"go1.12.5", Compiler:"gc", Platform:"linux/amd64"}
Server Version: version.Info{Major:"1", Minor:"15", GitVersion:"v1.15.0", GitCommit:"e8462b5b5d
c2584fdcd18e6bcfe9f1e4d970a529", GitTreeState:"clean", BuildDate:"2019-06-19T16:32:14Z", GoVers
ion:"go1.12.5", Compiler:"gc", Platform:"linux/amd64"}
$
```

Note: This says kubelet is version 1.15.2 and kubernetes (server) is v1.15.0

\$ kubectl get nodes

```
$ kubectl get nodes
NAME          STATUS    ROLES    AGE   VERSION
minikube      Ready     master   10m   v1.15.0
$
```

DEPLOYING AN APPLICATION TO A KUBERNETES CLUSTER requires a image to pull down and the number of replicas.

Before using **kubernetes create deployment** command, we will need to provide:

- the deployment name
- the app image location (including the full repository URL for images hosted outside of Docker hub)

\$ kubectl create deployment kubernetes-bootcamp --image=gcr.io/google-samples/kubernetes-bootcamp:v1

```
$ kubectl create deployment kubernetes-bootcamp --image=gcr.io/google-samples/kubernetes-bootca
mp:v1
deployment.apps/kubernetes-bootcamp created
$
```

What happened? Several steps occurred while creating your deployment:

1. Kubernetes (KDC) searched for a suitable node where the instance of the application could be run (which was trivial here since we only have one available node)
2. Kubernetes scheduled the application to run on the Node
3. Kubernetes configured the cluster to reschedule the instance on a new Node when needed

What's next? Confirmed your deployment using **\$ kubectl get deployments** to list your deployments

```
$ kubectl get deployments
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
kubernetes-bootcamp 1/1     1             1           4m10s
$
```

What's the result: You have a single deployment, called "kubernetes-bootcamp" and it is running a single instance of your app. The app instance is running inside a Docker container on your node.

View your app

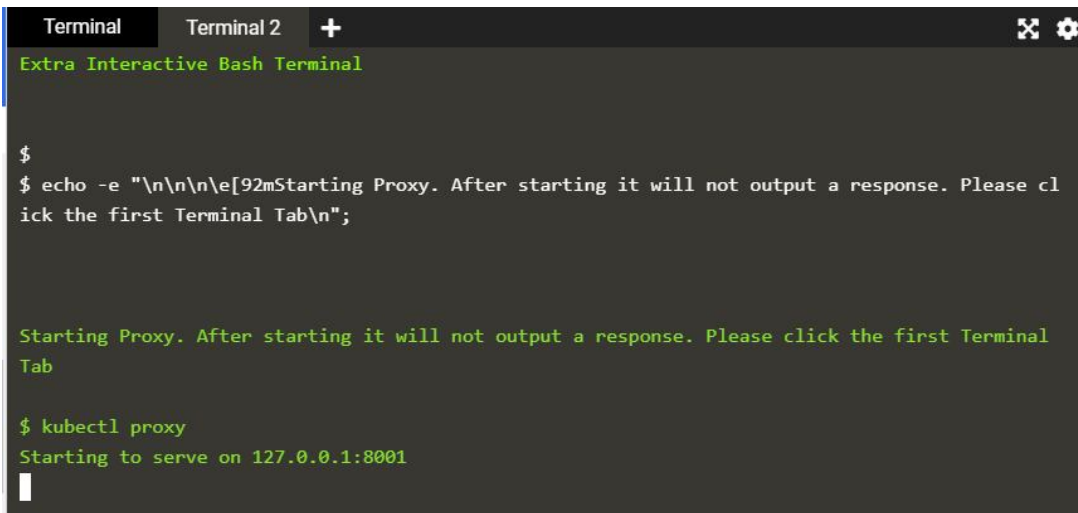
Actually as you know, your container is running within a Pod on inside of Kubernetes (in this case on the Node). The Pod is running on a private, isolated network. By default, they are visible to other pods and services within the same Kubernetes cluster, but not outside the network.

When we use kubectl, we are using an API endpoint to communicate to our application.

Using kubectl, we can create a Kube Proxy that will forward communications into the cluster-wide private network. The proxy can be terminated by pressing Ctrl-C and won't show any output while it is running.

LAB: Open a second terminal window to run the proxy: (use + to create 2nd terminal and enter supplied command)

```
echo -e "\n\n\n\e[92mStarting Proxy. After starting it will not output a response. Please click the first Terminal Tab\n"; kubectl proxy
```



```
Terminal  Terminal 2  +
Extra Interactive Bash Terminal

$
$ echo -e "\n\n\n\e[92mStarting Proxy. After starting it will not output a response. Please click the first Terminal Tab\n";
Starting Proxy. After starting it will not output a response. Please click the first Terminal Tab

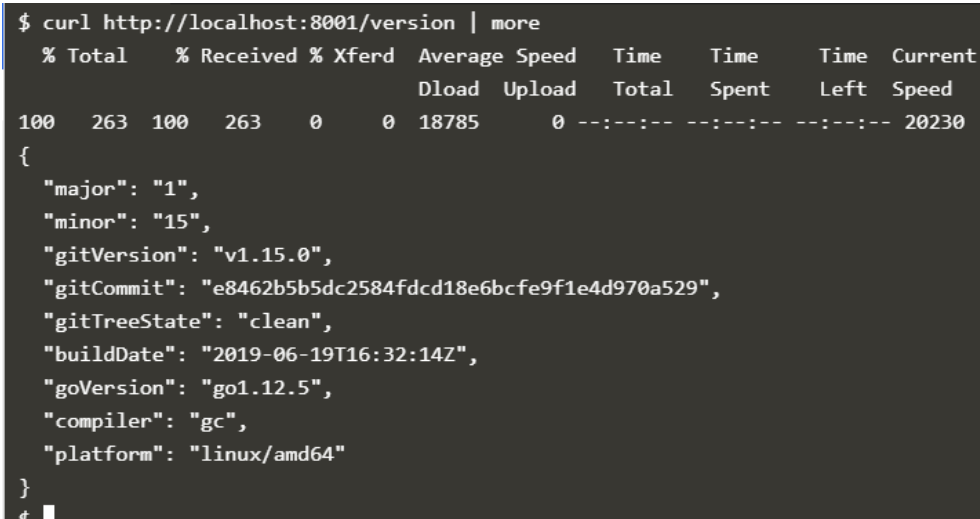
$ kubectl proxy
Starting to serve on 127.0.0.1:8001
```

Click on Terminal to continue:

Note: we now have a connection between our host (the online terminal) and the Kubernetes cluster. The Kube Proxy enables direct access to the API from the provided Lab Terminals.

NEXT

```
curl http://localhost:8001/version
```



```
$ curl http://localhost:8001/version | more
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           %             %             Dload  Upload  Total   Spent    Left     Speed
100    263    100    263     0     0  18785      0  --:--:-- --:--:-- --:--:-- 20230
{
  "major": "1",
  "minor": "15",
  "gitVersion": "v1.15.0",
  "gitCommit": "e8462b5b5dc2584fdcd18e6bcfe9f1e4d970a529",
  "gitTreeState": "clean",
  "buildDate": "2019-06-19T16:32:14Z",
  "goVersion": "go1.12.5",
  "compiler": "gc",
  "platform": "linux/amd64"
}
```

<https://kubernetes.io/docs/tutorials/kubernetes-basics/>

Since the Kube Proxy is running, the API Server will automatically create an endpoint for each pod, based on the pod name that is also accessible through the Proxy. You will need the Pod name, and we'll store it in the environment variable `POD_NAME`.

One way:

```
$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
kubernetes-bootcamp-75bccb7d87-xss9k 1/1     Running   0           27m
$
```

Lab way:

```
export POD_NAME=$(kubectl get pods -o go-template --template '{{.metadata.name}}{{"\n"}}{{end}}')
```

```
echo Name of the Pod: $POD_NAME
```

```
$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
kubernetes-bootcamp-75bccb7d87-xss9k 1/1     Running   0           27m
$ export POD_NAME=$(kubectl get pods -o go-template --template '{{range .items}}{{.metadata.name}}{{"\n"}}{{end}}')
$ echo Name of the Pod: $POD_NAME
Name of the Pod: kubernetes-bootcamp-75bccb7d87-xss9k
$
```

KEY FACT: `localhost:8001` requires the Kube Proxy to be running (as shown in Terminal 2). While the proxy still runs, our curl commands can work using `localhost:8001/path`.

KEY FINDING: In order for the new deployment to be accessible without using the Proxy, a **Service is required to be created. [Explained in future lab]**

First we need to get the Pod name, and we'll store it in the environment variable `POD_NAME`:

```
export POD_NAME=$(kubectl get
pods -o go-template --template
'{{range .items}}
{{.metadata.name}}{{"\n"}}
{{end}}')
echo Name of the Pod: $POD_NAME
↵
```

Note: Check the top of the terminal. The proxy was run in a new tab (Terminal 2), and the recent commands were executed the original tab (Terminal 1). The proxy still runs in the second tab, and this allowed our curl command to work using

`localhost:8001` .

Module 3 – Viewing Pods and Nodes; Troubleshooting deployed applications...

<https://kubernetes.io/docs/tutorials/kubernetes-basics/explore/explore-intro/>

A Pod is a group of one or more application containers (such as Docker or rkt) and includes shared storage (volumes), IP address and information about how to run them.

A “Pod” is a Kubernetes abstraction that represents a group of one or more application containers (like Docker or rkt) and some shared resources for those containers such as:

- shared storage or “Volumes”
- networking, such as a “unique cluster wide IP address”
- Information or metadata about each container, such as the container image version and specific port to use

Each Pod models an application-specific “logical host” and can contain different application containers that are relatively tightly coupled.

For instance a Pod might contain both the container with your Node.js app as well as a different container that feeds data to be published by the Node.js webserver. All containers within a single Pod share an IP address and common port space; are always co-located (same node) and co-scheduled (exist at same time); and run in the shared context on the same Worker Node.


Pods are the atomic unit on the Kubernetes platform. Pods are also referred to the unit of scheduling like a VM in VMware or a Process ID (pid) in Linux, Windows and UNIX.

When we create a Deployment on Kubernetes, that Deployment creates Pods with containers inside of them (as opposed to creating the containers directly).

Each Pod is tied to the Node where it is scheduled, and remains there until termination (according to the restart policy) or Pod deletion. In the case of a Node Failure, then identical Pods are scheduled on other available Nodes in the cluster.

Note: Other features of the cluster (Load balancing, etc) ensure that the externally available IP address doesn’t change and the Cluster will handle re-routing network traffic to the newly established Pod once it is available after Node Failure. In general cases, applications should have N+1 or N+2 instances (that translates to minimum of 3 or 4 instances) at all times for maximum balance of availability and maintainability.


Pods



Containers should only be scheduled together in a single Pod if they are tightly coupled and need to share resources such as disk. In the above diagrams, Pod1, Pod2 and Pod3 are more typical, with Pod 3 example having a dedicated helper node.

Note: A Volume exists or is deleted when the Pod terminates or is deleted. Learn more about “Persistent Volume” where are tied to the Cluster lifetime later to have more persistence across any given instance’s life cycle. External data stores can be arranged by the Kubernetes Administrator such as PostgreSQL, MemDB (temporary in-memory DB great for DevOps proof-of-concepts, etc), etc.

Node Overview



Nodes

A Pod always runs on a **Node**. The Node is a worker machine (i.e., “worker node”) in Kubernetes and may be either a virtual machine or a physical server, depending on the cluster. Each Node is managed by the master. A Node can have multiple pods, and the Kubernetes master automatically handles all scheduling of the pods across all available Nodes in the cluster, based on weighted formulas, available resources that match the deployment plan, etc. As stated, The Master’s automatic scheduling takes into account the available resources on each Node and maintains this metadata.

Every Kubernetes Node runs at least:

- Kubelet—a process responsible for communication between the Node and the Kubernetes Master; the process also manages the local Pods and containers running on a machine where it resides.
- A container runtime (like Docker, rkt), also known as container orchestration engine,--which is a container framework responsible for pulling the container image from a registry (repository), unpacking the container, and running the application in the Pod instance provided by Kubernetes.

Troubleshooting with Kubectl

In a previous lab, you starting using the kubectl command-line interface and interacted with the Master Node. In Module 3 you will contain to use ‘kubectl’ to get more infmromation about the Pods, the deployed applications and the environment that has been step.

In particular, you will use the most common operations or subcommands associated with kubectl, such as:

- **kubectl get** list cluster resources
- **kubectl describe** show more detailed information about a resource
- **kubectl logs** print the logs from a container within a pod
- **kubectl exec** execute commands directly on a container within a pod

Learn how to see when applications were deployed, what the applications’ current status are; where they are presently running in the Cluster and what their configuration are. Later with labels and other metadata, you will learn how to refine your commands to search for what you need with precision across larger and larger clusters with hundreds or thousands of nodes.

Given what we have discovered about our cluster components and the command line, start lab for Module 3.

Lab 3 “Exploring Your App” as companion to the Module 3 discussion

```
Terminal +
Kubernetes Bootcamp Terminal

$
$ sleep 1; launch.sh
Starting Kubernetes. This is expected to take less than a minute
Kubernetes Started
$ █
```

\$ kubectl get pods *Note: If you didn't see the single Pod, then please wait for Kubernetes cluster to start and try the command again*

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
kubernetes-bootcamp-5b48cfdcdb-vnnzx	1/1	Running	0	100s

As mentioned the next command **\$ kubectl describe pods** will get significantly more information.

```
$ kubectl describe pods
Name:          kubernetes-bootcamp-5b48cfdcdb-vnnzx
Namespace:     default
Priority:       0
Node:          minikube/172.17.0.13
Start Time:    Tue, 14 Jan 2020 03:32:33 +0000
Labels:        pod-template-hash=5b48cfdcdb
               run=kubernetes-bootcamp
Annotations:   <none>
Status:        Running
IP:            172.18.0.4
Controlled By: ReplicaSet/kubernetes-bootcamp-5b48cfdcdb
Containers:
  kubernetes-bootcamp:
    Container ID:  docker://67f0d0e6a435e5ff49790db9908a351e73ae75873d68fb326759b269a623efc1
    Image:         gcr.io/google-samples/kubernetes-bootcamp:v1
    Image ID:      docker-pullable://jocatalin/kubernetes-bootcamp@sha256:0d6b8ee63bb57c5f5b6156f446b3bc3b3c143d233037f3a2f00e279c8fcc64af
    Port:          8080/TCP
    Host Port:     0/TCP
    State:         Running
      Started:    Tue, 14 Jan 2020 03:32:34 +0000
Mounts:
  /var/run/secrets/kubernetes.io/serviceaccount from default-token-cb2mp (ro)
Conditions:
  Type           Status
  Initialized    True
  Ready          True
  ContainersReady True
  PodScheduled   True
Volumes:
  default-token-cb2mp:
    Type:          Secret (a volume populated by a Secret)
    SecretName:     default-token-cb2mp
    Optional:       false
QoS Class:        BestEffort
Node-Selectors:   <none>
Tolerations:      node.kubernetes.io/not-ready:NoExecute for 300s
                  node.kubernetes.io/unreachable:NoExecute for 300s
Events:
  Type     Reason      Age   From          Message
  ----     -
  Normal   Scheduled   3m34s default-scheduler Successfully assigned default/kubernetes-bootcamp-5b48cfdcdb-vnnzx to minikube
  Normal   Pulled      3m33s kubelet, minikube Container image "gcr.io/google-samples/kubernetes-bootcamp:v1" already present on machine
  Normal   Created     3m33s kubelet, minikube Created container kubernetes-bootcamp
  Normal   Started     3m33s kubelet, minikube Started container kubernetes-bootcamp
$ █
```

Using **\$ kubectl describe pod** we can see more details about the Pod's container such as:

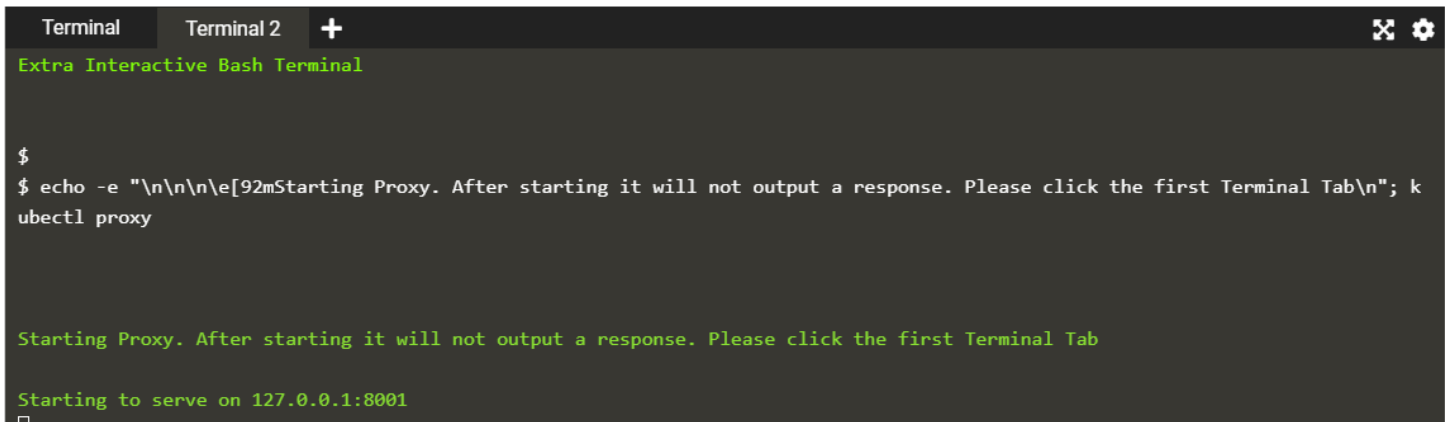
- IP address
- the ports in use and used
- a list of events related to the lifecycle of the Pod (get used to the order of events for alter troubleshooting)

In addition, you can use the **kubectl describe {other Kubernetes primitives}** to find out about most Kubernetes primitives such as nodes, pods, deployments, etc. The goal of **describe output** is it is designed to be **human readable**, **not to be scripted against** – which won't stop anyone from trying.

Lab 3.2 Show the App in the Terminal

As in the previous lab, we will use the "+" in the Session to create "Terminal 2" so we can establish a Kube Proxy somece the pods are running on an isolated, private network and we need to "proxy" access to them so we can interact and debug with them.

\$ echo -e "\n\n\n[e[92mStarting Proxy. After starting it will not output a response. Please click the first Terminat Tab\n"; kubectl proxy



```
Terminal  Terminal 2  +
Extra Interactive Bash Terminal

$
$ echo -e "\n\n\n[e[92mStarting Proxy. After starting it will not output a response. Please click the first Terminal Tab\n"; k
ubectl proxy

Starting Proxy. After starting it will not output a response. Please click the first Terminal Tab

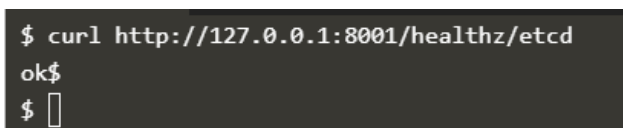
Starting to serve on 127.0.0.1:8001
```

As before click on "Terminal", and I like to verify you are able to access the resources; and that you can test the API:



```
Terminal  Terminal 2  +
$ kubectl get deployments
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
kubernetes-bootcamp 1/1     1            1           16m
$ kubectl get nodes
NAME     STATUS   ROLES    AGE   VERSION
minikube Ready    master   16m   v1.15.0
$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
kubernetes-bootcamp-5b48cfdcdb-vnnzx 1/1     Running   0          16m
$
```

Here is one of my simple API Test commands after some exploring after Module 2 (see Appendix)



```
$ curl http://127.0.0.1:8001/healthz/etcd
ok$
$
```

Lab 3.2 continues

Now again, we'll get the Pod name and query that pod directly through the proxy. To get the Pod name and store it in the `POD_NAME` environment variable:

```
export POD_NAME=$(kubectl get pods -o go-template --template '{{range .items}}{{.metadata.name}}{{"\n"}}{{end}}') echo Name of the Pod: $POD_NAME
```

Note: you can see that the “Name of the Pod” output is same as what kubectl get pods returned above.

```
$ export POD_NAME=$(kubectl get pods -o go-template --template '{{range .items}}{{.metadata.name}}{{"\n"}}{{end}}')
$ echo Name of the Pod: $POD_NAME
Name of the Pod: kubernetes-bootcamp-5b48cfdbc-d-vnnzx
$
```

To see the output of our application, run a curl request.

```
curl http://localhost:8001/api/v1/namespaces/default/pods/$POD_NAME/proxy/
```

The url is the route to the API of the Pod. Here is the output:

```
$ curl http://localhost:8001/api/v1/namespaces/default/pods/$POD_NAME/proxy/
Hello Kubernetes bootcamp! | Running on: kubernetes-bootcamp-5b48cfdbc-d-vnnzx | v=1
$
```

Lab 3.3 View the Container Logs

Since the prior lab created an environment variable, we can simply use `$POD_NAME` instead of dealing with the longer name of the pod and avoid mistakes. `$ kubectl logs $POD_NAME` will provide the following output:

```
$ kubectl logs $POD_NAME
Kubernetes Bootcamp App Started At: 2020-01-14T03:32:34.740Z | Running On: kubernetes-bootcamp-5b48cfdbc-d-vnnzx

Running On: kubernetes-bootcamp-5b48cfdbc-d-vnnzx | Total Requests: 1 | App Uptime: 1361.899 seconds | Log Time: 2020-01-14T03:55:16.639Z
$
```

Lab 3.4 Executing command on the actual Container

Since the Pod is up and running, you can address commands, using `kubectl exec {pod-name} {command}` as follows:

`$ kubectl exec $POD_NAME env`

```
$ kubectl exec $POD_NAME env
PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
HOSTNAME=kubernetes-bootcamp-5b48cfdbc-d-vnnzx
KUBERNETES_PORT_443_TCP_PORT=443
KUBERNETES_PORT_443_TCP_ADDR=10.96.0.1
KUBERNETES_SERVICE_HOST=10.96.0.1
KUBERNETES_SERVICE_PORT=443
KUBERNETES_SERVICE_PORT_HTTPS=443
KUBERNETES_PORT=tcp://10.96.0.1:443
KUBERNETES_PORT_443_TCP=tcp://10.96.0.1:443
KUBERNETES_PORT_443_TCP_PROTO=tcp
NPM_CONFIG_LOGLEVEL=info
NODE_VERSION=6.3.1
HOME=/root
$
```

<https://kubernetes.io/docs/tutorials/kubernetes-basics/>

Note: It is good habit to specify the POD NAME even though, with a single Pod on a single Node it could be omitted.

Creating a Bash Session on the Pod's container:

\$ kubectl exec -ti \$POD_NAME bash

```
$ kubectl exec -ti $POD_NAME bash
root@kubernetes-bootcamp-5b48cfdbc-d-vnnzx:/# wc -l server.js
19 server.js
root@kubernetes-bootcamp-5b48cfdbc-d-vnnzx:/#
```

Since we now have an open console on the container hosting our NodeJS application, we can show the source code of our app located in server.js file. But first, I was curious on how many lines of code are in this app, server.js

\$ wc -l server.js above is a simple Linux command that Ubuntu supports that says the app is 19 lines long.

However, the lab wants you to run **/# cat server.js** to see the entire app in NodeJS container

```
root@kubernetes-bootcamp-5b48cfdbc-d-vnnzx:/# cat server.js
var http = require('http');
var requests=0;
var podname= process.env.HOSTNAME;
var startTime;
var host;
var handleRequest = function(request, response) {
  response.setHeader('Content-Type', 'text/plain');
  response.writeHead(200);
  response.write("Hello Kubernetes bootcamp! | Running on: ");
  response.write(host);
  response.end(" | v=1\n");
  console.log("Running On:" ,host, "| Total Requests:", ++requests,"| App Uptime:", (new Date() - startTime)/1000 , "seconds", "| Log Time:",new Date());
}
var www = http.createServer(handleRequest);
www.listen(8080,function () {
  startTime = new Date();
  host = process.env.HOSTNAME;
  console.log ("Kubernetes Bootcamp App Started At:",startTime, "| Running On: " ,host, "\n" );
});
root@kubernetes-bootcamp-5b48cfdbc-d-vnnzx:/#
```

Now let's confirm the application, server.js, is up and running by testing the API again:

```
root@kubernetes-bootcamp-5b48cfdbc-d-vnnzx:/# curl localhost:8080
Hello Kubernetes bootcamp! | Running on: kubernetes-bootcamp-5b48cfdbc-d-vnnzx | v=1
root@kubernetes-bootcamp-5b48cfdbc-d-vnnzx:/#
```

Note: here we used localhost because we executed the command inside the NodeJS Pod. If you cannot connect to localhost:8080, check to make sure you have run the kubectl exec command and are launching the command from within the Pod.

After typing exit, you will leave the connection to Nodejs container and can review the logs for the Pod again:

```
$ kubectl exec -ti $POD_NAME bash
root@kubernetes-bootcamp-5b48cfdbc-d-vnnzx:/# ls
bin boot core dev etc home lib lib64 media mnt opt proc root run sbin server.js srv sys tmp usr var
root@kubernetes-bootcamp-5b48cfdbc-d-vnnzx:/# exit
exit
$
```

You can see the logs below:

```
$ kubectl logs $POD_NAME
Kubernetes Bootcamp App Started At: 2020-01-14T03:32:34.740Z | Running On: kubernetes-bootcamp-5b48cfdbc-d-vnnzx

Running On: kubernetes-bootcamp-5b48cfdbc-d-vnnzx | Total Requests: 1 | App Uptime: 1361.899 seconds | Log Time: 2020-01-14T03:55:16.639Z
Running On: kubernetes-bootcamp-5b48cfdbc-d-vnnzx | Total Requests: 2 | App Uptime: 2234.447 seconds | Log Time: 2020-01-14T04:09:49.187Z
$
```

Module 4 – Using a Service to Expose your App PUBLICLY

<https://kubernetes.io/docs/tutorials/kubernetes-basics/explore/explore-intro/>

All Kubernetes Pods have a lifecycle. When a worker node dies, the Pods running on the Node are also lost. Any Volumes attached to the Pod are lost. The state is gone.

With a failed Node, a ReplicaSet might then dynamically drive the cluster back to the desired state via creation of new Pods identical to the lost Pods to keep your application running for the next set of input received.

Another example might be an image-processing backend with 3 replicas. Those replicas are exchangeable; the front-end system should not care about backend replicas or even if a Pod is lost and recreated.

Summary – each Pod in a Kubernetes cluster has a unique IP address, even Pods on the same Node, so there needs to be a way of automatically reconciling changes among Pods so that your applications can continue to function.

*A Pod is a group of one or more application containers (such as Docker or rkt) and includes shared storage (volumes), IP address and information about how to run them. The previous lab used Kube Proxy to provide a way of reaching into the unique private address space assigned to a given Pod and we tested the access using curl command. However Kubernetes provides the means to expose an application outside the Kubernetes cluster using a **SERVICE**.*

A Kubernetes Service is an abstraction layer which defines a logical set of Pods and enables external traffic exposure, load balancing and service discovery for those Pods.

A “SERVICE” in Kubernetes is an abstraction which defines a **logical set of Pods and a policy by which to access them**. Services enable a *loose coupling* between dependent Pods. **A Service is defined using YAML (preferred, or JSON as last resort)**, like all Kubernetes objects. The set of Pods targeted by a Service is usually determined by a **LabelSelector** (see below why you might want a Service without including **selector** in the spec).

A SERVICE ALLOWS YOUR APPLICATIONS TO RECEIVE TRAFFIC. Without a Service, even though each Pod has a unique IP address, it is private to the cluster, and those IPs are not exposed outside of the cluster without a Service. Services can be exposed in different ways by specifying a **type** in the **ServiceSpec**:

- **ClusterIP (default)**—exposes the service on the internal IP in the cluster. This type makes the service only reachable from within the cluster itself.
- **NodePort**—exposes the Service on the same port of each selected Node in the cluster using NAT. Makes a Service accessible from outside the cluster using <NodeIP>:<NodePort>. This is a SuperSet of ClusterIP.
- **LoadBalancer**—creates an external load balancer in the current cloud (if supported) and assigns a fixed, external IP to the Service. Superset of NodePort.
- **ExternalName**—exposes the Service using an arbitrary name (specified by **externalName** in the spec) by returning a CNAME record with the name. No proxy is used. This type requires Kubernetes v1.7 or higher of kube-dns.

Refer to “Using Source IP” (<https://kubernetes.io/docs/tutorials/services/source-ip/>) tutorial for more details on the different types of Services. Also see “Connecting Applications with Services” (<https://kubernetes.io/docs/concepts/services-networking/connect-applications-service>).

Note: A service created without **selector** will also not create the corresponding Endpoints object. This allows users to manually map a Service to specific endpoints. Another possibility is that you are strictly using **type: ExternalName** and there is no selector.

Services and Labels

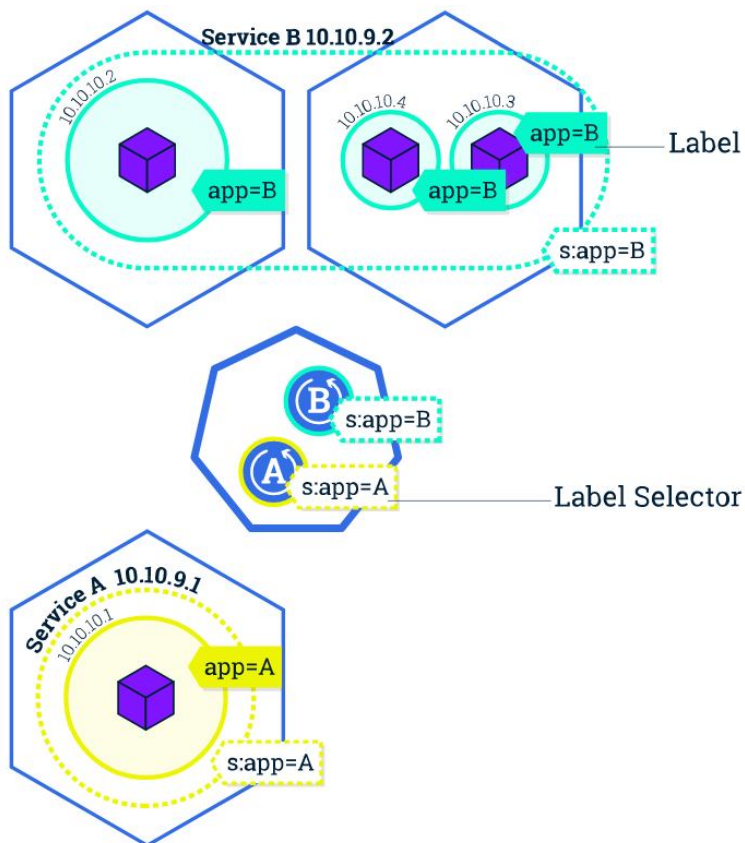

Did you know...

You can create a Service at the same time you create a Deployment by using `--expose` in `kubectl`.

The role of a Service is to route network traffic across a set of Pods. Services are abstractions that allow pods to die and replicate in Kubernetes without impacting your application. Discovery and routing among dependent Pods (such as the front-end and back-end components in an application) is handled by Kubernetes Services.

Services match a set of Pods using **labels** () and **selectors** (), which is a grouping primitive that allows logical operations on objects in Kubernetes. Labels are key/value pairs attached to objects and can be used in many ways:

- designate objects for development, test, build and production
- embed version tags
- classify an object using tags



In summary, Labels can be attached to objects at creation time or can be added later on.

Labels can be modified at any time.

The Lab 4 for Module 4 will expose our application publicly using a Service and then spend some time applying labels.

Lab 4 “Exposing Your App Publicly using a Service” as companion to the Module 4 discussion

```
Terminal +
Kubernetes Bootcamp Terminal

$
$ sleep 1; launch.sh
Starting Kubernetes...
█
```

```
Terminal +
Kubernetes Bootcamp Terminal

$
$ sleep 1; launch.sh
Starting Kubernetes...
Kubernetes Started
$ █
```

Some times you have to wait for your provisioning (45 seconds...) before you see “Kubernetes Started”. Now, run a few commands to know the state of your cluster that took time to start up:

Suggested: \$ kubectl get pods
 \$ kubectl get deployments
 \$ kubectl get services

\$ kubectl get pods *Note: If you didn't see the single Pod, then please wait for Kubernetes cluster to start and try again*

```
$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
kubernetes-bootcamp-5c69669756-rk8zp 1/1     Running   0          1m

$ kubectl get deployments
NAME                DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
kubernetes-bootcamp 1          1         1             1           1m

$ kubectl get services
NAME         TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)    AGE
kubernetes  ClusterIP   10.96.0.1    <none>        443/TCP    1m
$ █
```

Currently, we are running a Service called kubernetes-bootcamp. The Service received a unique cluster-IP, an internal port and an external IP (the IP of the Node). – however, we must use “expose” option for kubectl

```
$ kubectl expose deployment/kubernetes-bootcamp --type="NodePort" --port 8080
service/kubernetes-bootcamp exposed

$ kubectl get services
NAME                TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)            AGE
kubernetes          ClusterIP   10.96.0.1    <none>        443/TCP            10m
kubernetes-bootcamp NodePort     10.100.94.226 <none>        8080:31991/TCP     6s
$ █
```

To find out what port was opened externally (by the NodePort option), please run the **describe service** command

```
$ kubectl describe services/kubernetes-bootcamp
Name:                kubernetes-bootcamp
Namespace:           default
Labels:              run=kubernetes-bootcamp
Annotations:         <none>
Selector:            run=kubernetes-bootcamp
Type:                NodePort
IP:                  10.100.94.226
Port:                <unset> 8080/TCP
TargetPort:          8080/TCP
NodePort:            <unset> 31991/TCP
Endpoints:           172.18.0.4:8080
Session Affinity:    None
$ █
```

<https://kubernetes.io/docs/tutorials/kubernetes-basics/>

Create an environment variable called `NODE_PORT` that has the value of the Node port assigned:

```
export NODE_PORT=$(kubectl get services/kubernetes-bootcamp -o go-template='{{(index .spec.ports 0).nodePort}}') echo NODE_PORT=$NODE_PORT
```

Results:

```
Terminal +
$ export NODE_PORT=$(kubectl get services/kubernetes-bootcamp -o go-template='{{(index .spec.ports 0).nodePort}}')
$ echo NODE_PORT=$NODE_PORT
NODE_PORT=31991
$
```

Now, test that the app is exposed outside of the cluster using `curl`, the IP of the Node and the externally exposed port:

```
curl $(minikube ip):$NODE_PORT
```

And we get a response from the server. The Service is exposed.

```
Terminal +
$ export NODE_PORT=$(kubectl get services/kubernetes-bootcamp -o go-template='{{(index .spec.ports 0).nodePort}}')
$ echo NODE_PORT=$NODE_PORT
NODE_PORT=31991
$ curl $(minikube ip):$NODE_PORT
Hello Kubernetes bootcamp! | Running on: kubernetes-bootcamp-5c69669756-rk8zp | v=1
$
```

Lab 4.2 Using Labels (as part of Exposing your App Publicly with Services)

So far, exposed our service using the `kubectl expose` command, we were able to describe the service and identify the `NodePort` assigned and put that information into an environment variable `$NODE_PORT` we subsequently used to use the `curl $(minikube ip):$NODE_PORT` command to test that the app is actually exposed outside of the cluster:

```
Terminal +
$ export NODE_PORT=$(kubectl get services/kubernetes-bootcamp -o go-template='{{(index .spec.ports 0).nodePort}}')
$ echo NODE_PORT=$NODE_PORT
NODE_PORT=31991
$ curl $(minikube ip):$NODE_PORT
Hello Kubernetes bootcamp! | Running on: kubernetes-bootcamp-5c69669756-rk8zp | v=1
$
```

Lab 4.2 continues

In this lab, 4.2, we are going to use labels – and to do that, we have to learn the automatically created label the Deployment assigned to our Pod. We do this with the **\$ kubectl describe deployment**

```
$ kubectl describe deployment | more
Name:                kubernetes-bootcamp
Namespace:           default
CreationTimestamp:    Tue, 14 Jan 2020 04:57:37 +0000
Labels:              run=kubernetes-bootcamp
Annotations:         deployment.kubernetes.io/revision=1
Selector:            run=kubernetes-bootcamp
Replicas:            1 desired | 1 updated | 1 total | 1 available | 0 unavailable
StrategyType:        RollingUpdate
MinReadySeconds:     0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels:  run=kubernetes-bootcamp
  Containers:
    kubernetes-bootcamp:
      Image:        gcr.io/google-samples/kubernetes-bootcamp:v1
      Port:         8080/TCP
      Host Port:    0/TCP
      Environment:  <none>
      Mounts:       <none>
      Volumes:      <none>
  Conditions:
    Type           Status  Reason
    ----           -
    Available       True    MinimumReplicasAvailable
    Progressing     True    NewReplicaSetAvailable
OldReplicaSets:    <none>
NewReplicaSet:     kubernetes-bootcamp-5c69669756 (1/1 replicas created)
Events:
  Type           Reason              Age   From                      Message
  ----           -
  Normal         ScalingReplicaSet   27m   deployment-controller     Scaled up replica set kubernetes-bootcamp-5c69669756 to 1
$
```

The fourth line provides the value “run=kubernetes-bootcamp” as the label

```
Labels: run=kubernetes-bootcamp
```

Use this value with the `-l` option on “kubectl get pods” and “kubectl get services” as follows:

```
$ kubectl get pods -l run=kubernetes-bootcamp
NAME                                READY   STATUS    RESTARTS   AGE
kubernetes-bootcamp-5c69669756-rk8zp 1/1     Running   0           31m
$ kubectl get services -l run=kubernetes-bootcamp
NAME           TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)          AGE
kubernetes-bootcamp NodePort    10.100.94.226 <none>        8080:31991/TCP   20m
```

Once again, put the name of this Pod in `$POD_NAME`:

```
$ export POD_NAME=$(kubectl get pods -o go-template --template '{{range .items}}{{.metadata.name}}{{"\n"}}{{end}}')
$ echo Name of the Pod: $POD_NAME
Name of the Pod: kubernetes-bootcamp-5c69669756-rk8zp
$
```

<https://kubernetes.io/docs/tutorials/kubernetes-basics/>

To apply a new label we use the label command followed by the object type, object name and the new label:

```
kubectl label pod $POD_NAME app=v1
```

```
$ kubectl label pod $POD_NAME app=v1
pod/kubernetes-bootcamp-5c69669756-rk8zp labeled
```

This will apply a new label to our Pod (we pinned the application version to the Pod), and we can check it with the describe pod command: `kubectl describe pods $POD_NAME`

```
$ kubectl describe pods $POD_NAME | more
Name:          kubernetes-bootcamp-5c69669756-rk8zp
Namespace:     default
Node:          minikube/172.17.0.26
Start Time:    Tue, 14 Jan 2020 04:57:55 +0000
Labels:        app=v1
               pod-template-hash=1725225312
               run=kubernetes-bootcamp
Annotations:   <none>
Status:        Running
IP:            172.18.0.4
Controlled By: ReplicaSet/kubernetes-bootcamp-5c69669756
Containers:
  kubernetes-bootcamp:
    Container ID:  docker://87498c04a988734a0a1794e2863dfe4db86f6f160f910535546325133eec5955
    Image:         gcr.io/google-samples/kubernetes-bootcamp:v1
    Image ID:      docker-pullable://gcr.io/google-samples/kubernetes-bootcamp@sha256:0d6b8ee63bb57c5f5b6156f446b3bc3b3c143d233037f3a2f00e279c8fcc64af
    Port:          8080/TCP
    Host Port:     0/TCP
    State:         Running
      Started:     Tue, 14 Jan 2020 04:57:56 +0000
    Ready:         True
    Restart Count: 0
    Environment:   <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-jdzdx (ro)
Conditions:
  Type            Status
  Initialized     True
  Ready           True
  PodScheduled    True
Volumes:
  default-token-jdzdx:
    Type:          Secret (a volume populated by a Secret)
    SecretName:    default-token-jdzdx
    Optional:      false
QoS Class:        BestEffort
Node-Selectors:   <none>
Tolerations:      node.kubernetes.io/not-ready:NoExecute for 300s
                  node.kubernetes.io/unreachable:NoExecute for 300s
Events:
  Type    Reason              Age             From              Message
  ----    -
  Warning FailedScheduling    35m (x5 over 35m) default-scheduler 0/1 nodes are available: 1 node(s) were not ready.
  Normal  Scheduled           35m             default-scheduler Successfully assigned kubernetes-bootcamp-5c69669756-rk8zp to minikube
  Normal  SuccessfulMountVolume 35m             kubelet, minikube MountVolume.SetUp succeeded for volume "default-token-jdzdx"
  Normal  Pulled               35m             kubelet, minikube Container image "gcr.io/google-samples/kubernetes-bootcamp:v1" already present on machine
  Normal  Created              35m             kubelet, minikube Created container
  Normal  Started              35m             kubelet, minikube Started container
$
```

We see here that the label is attached now to our Pod. And we can query now the list of pods using the new label:

```
kubectl get pods -l app=v1
```

And we see the Pod.

```
$ kubectl get pods -l app=v1
```

NAME	READY	STATUS	RESTARTS	AGE
kubernetes-bootcamp-5c69669756-rk8zp	1/1	Running	0	37m

```
$
```

Lab 4.3 Deleting a Service (as part of Exposing your App Publicly with Services)

Kubectl has a “delete service” option that will allow you to clean up. During the test, you should validate what you intended did occur:

```
$ kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	39m
kubernetes-bootcamp	NodePort	10.100.94.226	<none>	8080:31991/TCP	28m

```
$ kubectl delete service -l run=kubernetes-bootcamp
```

service "kubernetes-bootcamp" deleted

```
$ kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	39m

```
$
```

You can further demonstrate the Service was removed using the `$ curl $(minikube ip):$NODE_PORT` command again:

```
$ curl $(minikube ip):$NODE_PORT
```

curl: (7) Failed to connect to 172.17.0.26 port 31991: Connection refused

```
$
```

Last, is your app still running? You can verify in one command by having the Container within the Pod execute it:

```
$ kubectl exec -ti $POD_NAME curl localhost:8080
```

```
$ kubectl exec -ti $POD_NAME curl localhost:8080
```

Hello Kubernetes bootcamp! | Running on: kubernetes-bootcamp-5c69669756-rk8zp | v=1

```
$
```

So we see the application is up and running. Remember, the deployment is managing the application. To complete the lifecycle, you will have to shut down the application which means you need to delete the Deployment as well.

Module 5 – Running Multiple Instances of Your App

Objectives: **Scale an app using kubectl.**


<https://kubernetes.io/docs/tutorials/kubernetes-basics/scale/scale-intro/>

Scaling an application

In the previous modules we created a [Deployment](#), and then exposed it publicly via a [Service](#). The Deployment created only one Pod for running our application. When traffic increases, we will need to scale the application to keep up with user demand. **Scaling** is accomplished by **changing the number of replicas in a Deployment**

You can scale from the start, using the `--replica` parameter for the `kubectl run` command, when you create a Deployment with multiple instances

Scaling overview



Scaling out a Deployment will ensure new Pods are created and scheduled to Nodes with available resources. Scaling will increase the number of Pods to the new desired state. Kubernetes also supports [autoscaling](#) of Pods, but it is outside of the scope of this tutorial. Scaling to zero is also possible, and it will terminate all Pods of the specified Deployment.

Running multiple instances of an application will require a way to distribute the traffic to all of them. Services have an integrated load-balancer that will distribute network traffic to all Pods of an exposed Deployment. Services will monitor continuously the running Pods using endpoints, to ensure the traffic is sent only to available Pods.

Once you have multiple instances of an Application running, you would be able to do Rolling updates without downtime. We'll cover that in the next module. Now, let's go to the online terminal and scale our application.

➔ *Scaling is accomplished by changing the number of replicas in a Deployment.*

Lab 5 “Scaling Your App” as companion to the Module 5 discussion

The goal of this interactive scenario is to scale a deployment with `kubectl scale` and to see the load balancing in action

The online terminal is a pre-configured Linux environment that can be used as a regular console (you can type commands). Clicking on the blocks of code followed by the ENTER sign will execute that command in the terminal.

```
Kubernetes Bootcamp Terminal
$
$ sleep 1; launch.sh
Starting Kubernetes. This is expected to take less than a minute
Kubernetes Started
$
```

Some times you have to wait for your provisioning (45 seconds...) before you see “**Kubernetes Started**”. Now, run a few commands to know the state of your cluster that took time to start up:

Suggested:

- \$ `kubectl get pods`
- \$ `kubectl get deployments`
- \$ `kubectl get services`

```
Kubernetes Bootcamp Terminal
$
$ sleep 1; launch.sh
Starting Kubernetes. This is expected to take less than a minute
Kubernetes Started
$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
kubernetes-bootcamp-5b48cfdbc-dhjm8 1/1     Running   0           61s
$ kubectl get nodes
NAME      STATUS   ROLES    AGE   VERSION
minikube  Ready    master   76s   v1.15.0
$ kubectl get deployments
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
kubernetes-bootcamp  1/1     1             1           82s
$ kubectl get services
NAME          TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)          AGE
kubernetes    ClusterIP   10.96.0.1    <none>        443/TCP          88s
kubernetes-bootcamp  NodePort    10.111.171.95 <none>        8080:32466/TCP   86s
$
```

Lab for Step 5.1: Scaling a deployment

Confirm you deployments by using the \$ `kubectl get deployments` (see output above). This shows we have once 1 Pod with the Ready column sure the ratio of **Current** to **Desired** replicas. It is important in large environments to monitor that **CURRENT** is the number of replicas running now, and **DESIRED** will be the configured of number of replicas that should be running. Likewise, **UP-TO-DATE** is the number of replicas updated to match the desired (configured) state; and **AVAILABLE** state shows how many replicas are actually AVAILABLE to the users.

Lab for Lab 5.1 continues

Now, run the following commands:

```
$ kubectl get deployments
$ kubectl scale deployments/kubernetes-bootcamp --replicas=4
$ kubectl get deployments
$ kubectl get pods --o wide
```

Here is a perfect result – you can see the interim stages

- READY = 1/4 right after the scale command; then
- 2 Pods are ready (1/1) with the Status of Running, with an IPv4 address each; but the other
- 2 Pods are not ready (0/1) with the Status of ContainerCreating, and <none> as the IPv4 value
- Finally a subsequent **\$ kubectl get deployments** shows READY is 4/4, Update is 4 and AVAILABLE is 4

```
NAME                READY    UP-TO-DATE    AVAILABLE    AGE
kubernetes-bootcamp 1/1      1             1            9m9s
$ kubectl scale deployments/kubernetes-bootcamp --replicas=4
deployment.extensions/kubernetes-bootcamp scaled
$ kubectl get deployments
NAME                READY    UP-TO-DATE    AVAILABLE    AGE
kubernetes-bootcamp 1/4      4             1            9m12s
$ kubectl get pods -o wide
NAME                READY    STATUS              RESTARTS   AGE    IP             NODE       NOMINATED NODE
READINESS GATES
kubernetes-bootcamp-5b48cfdbc-d79n7x 1/1      Running           0          2s     172.18.0.6     minikube   <none>
kubernetes-bootcamp-5b48cfdbc-dhjm8 1/1      Running           0          9m7s   172.18.0.4     minikube   <none>
kubernetes-bootcamp-5b48cfdbc-hpmfb 0/1      ContainerCreating  0          2s     <none>         minikube   <none>
kubernetes-bootcamp-5b48cfdbc-vslzm 0/1      ContainerCreating  0          2s     <none>         minikube   <none>
$ kubectl get deployments
NAME                READY    UP-TO-DATE    AVAILABLE    AGE
kubernetes-bootcamp 4/4      4             4            9m17s
```

For good measure, I ran the **\$ kubectl get pods --o wide** which shows the correct READY ratio and a valid IPv4 address filled in

```
$ kubectl get pods -o wide
NAME                READY    STATUS              RESTARTS   AGE    IP             NODE       NOMINATED NODE   READINESS GATES
kubernetes-bootcamp-5b48cfdbc-d79n7x 1/1      Running           0          5m23s   172.18.0.6     minikube   <none>            <none>
kubernetes-bootcamp-5b48cfdbc-dhjm8 1/1      Running           0          14m     172.18.0.4     minikube   <none>            <none>
kubernetes-bootcamp-5b48cfdbc-hpmfb 1/1      Running           0          5m23s   172.18.0.7     minikube   <none>            <none>
kubernetes-bootcamp-5b48cfdbc-vslzm 1/1      Running           0          5m23s   172.18.0.8     minikube   <none>            <none>
```

There are 4 Pods now, with different IP addresses. The change was registered in the Deployment events log. To check that, use the describe command:

```
kubectl describe deployments/kubernetes-bootcamp
```



```
$ kubectl describe deployments/kubernetes-bootcamp | more
Name:                kubernetes-bootcamp
Namespace:           default
CreationTimestamp:    Fri, 17 Jan 2020 04:56:11 +0000
Labels:              run=kubernetes-bootcamp
Annotations:         deployment.kubernetes.io/revision: 1
Selector:            run=kubernetes-bootcamp
Replicas:            4 desired | 4 updated | 4 total | 4 available | 0 unavailable
StrategyType:        RollingUpdate
MinReadySeconds:      0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels:  run=kubernetes-bootcamp
  Containers:
    kubernetes-bootcamp:
      Image:        gcr.io/google-samples/kubernetes-bootcamp:v1
      Port:         8080/TCP
      Host Port:    0/TCP
      Environment:  <none>
      Mounts:       <none>
      Volumes:      <none>
  Conditions:
    Type           Status  Reason
    ----           -
    Progressing    True    NewReplicaSetAvailable
    Available      True    MinimumReplicasAvailable
OldReplicaSets:  <none>
NewReplicaSet:   kubernetes-bootcamp-5b48cfdcdb (4/4 replicas created)
Events:
  Type     Reason             Age   From                  Message
  ----     -
  Normal   ScalingReplicaSet  17m   deployment-controller  Scaled up replica set kubernetes-bootcamp-5b48cfdcdb to 1
  Normal   ScalingReplicaSet  8m7s  deployment-controller  Scaled up replica set kubernetes-bootcamp-5b48cfdcdb to 4
$
```

You can also view in the output of this command that there are 4 replicas now.

```
Replicas: 4 desired | 4 updated | 4 total | 4 available | 0 unavailable
```

Lab for Step 5.2: Load Balancing

Let's verify whether the Service is load-balancing the traffic like it showed for Service A in the Module 5 diagram. Once again, use the same describe command, but let's look for the "exposed IP" and "exposed Port", but this time use `services/kubernetes-bootcamp` instead of `deployments/kubernetes-bootcamp`.

```
kubectl describe services/kubernetes-bootcamp
```

```
$ kubectl get services
NAME                TYPE        CLUSTER-IP    EXTERNAL-IP    PORT(S)        AGE
kubernetes          ClusterIP   10.96.0.1     <none>         443/TCP        23m
kubernetes-bootcamp NodePort    10.111.171.95 <none>         8080:32466/TCP 23m
$ kubectl describe services/kubernetes-bootcamp
Name:                kubernetes-bootcamp
Namespace:           default
Labels:              run=kubernetes-bootcamp
Annotations:         <none>
Selector:            run=kubernetes-bootcamp
Type:               NodePort
IP:                 10.111.171.95
Port:               <unset> 8080/TCP
TargetPort:         8080/TCP
NodePort:           <unset> 32466/TCP
Endpoints:          172.18.0.4:8080,172.18.0.6:8080,172.18.0.7:8080 + 1 more...
Session Affinity:    None
External Traffic Policy: Cluster
Events:             <none>
$
```

As before, let's create an environment variable called `NODE_PORT` for the value of `NODE_PORT` above so we can use `curl` to confirm the exposed IP and port is working (run `curl` command multiple times)

```
export NODE_PORT=$(kubectl get services/kubernetes-bootcamp -o go-template='{{(index .spec.ports 0).nodePort}}')
echo NODE_PORT=$NODE_PORT
```

```
$ export NODE_PORT=$(kubectl get services/kubernetes-bootcamp -o go-template='{{(index .spec.ports 0).nodePort}}')
$ echo NODE_PORT=$NODE_PORT
NODE_PORT=32466
$ curl $(minikube ip):$NODE_PORT
Hello Kubernetes bootcamp! | Running on: kubernetes-bootcamp-5b48cfdcdb-79n7x | v=1
$ curl $(minikube ip):$NODE_PORT
Hello Kubernetes bootcamp! | Running on: kubernetes-bootcamp-5b48cfdcdb-vs1zm | v=1
$ curl $(minikube ip):$NODE_PORT
Hello Kubernetes bootcamp! | Running on: kubernetes-bootcamp-5b48cfdcdb-vs1zm | v=1
$ curl $(minikube ip):$NODE_PORT
Hello Kubernetes bootcamp! | Running on: kubernetes-bootcamp-5b48cfdcdb-hpmfb | v=1
$ curl $(minikube ip):$NODE_PORT
Hello Kubernetes bootcamp! | Running on: kubernetes-bootcamp-5b48cfdcdb-hpmfb | v=1
$ curl $(minikube ip):$NODE_PORT
Hello Kubernetes bootcamp! | Running on: kubernetes-bootcamp-5b48cfdcdb-hpmfb | v=1
```

We can see, we have 3 different Pods values for the six times we ran the `$ curl $(minikube ip):$NODE_PORT` command. It isn't exactly round-robin (79n7x = 16%, vs1zm = 34% and hpmfb = 50%), but it is load balancing.

Lab for Step 5.3: Scale Down

While the peaks and valleys of network traffic from the Internet will normally shape your Kubernetes deployment, we can demonstrate by manually **scaling down the cluster**.

Let's alter the cluster by reducing the deployment by 50 percent (going from 4 replicas to 2 replicas).

You accomplish this by using the same exact SCALE parameter, but use **replicas=2** instead of **replicas=4**

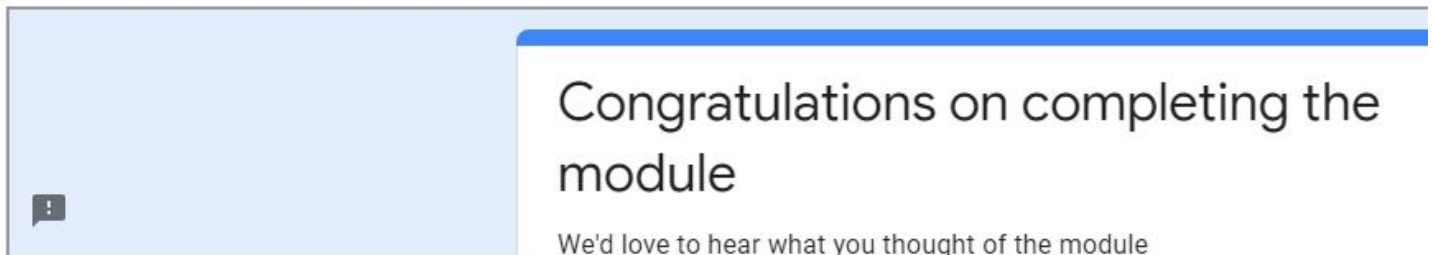
```
$ kubectl scale deployments/kubernetes-bootcamp --replicas=2
```

```
$ kubectl get deployments
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
kubernetes-bootcamp 4/4      4             4           36m
$ kubectl scale deployments/kubernetes-bootcamp --replicas=2
deployment.extensions/kubernetes-bootcamp scaled
$ kubectl get deployments
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
kubernetes-bootcamp 2/2      2             2           36m
$ kubectl get pods -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP            NODE       NOMINATED NODE   READINESS GATES
kubernetes-bootcamp-5b48cfdbcdb-79n7x 1/1     Running   0          27m   172.18.0.6    minikube   <none>            <none>
kubernetes-bootcamp-5b48cfdbcdb-dhjm8 1/1     Running   0          36m   172.18.0.4    minikube   <none>            <none>
kubernetes-bootcamp-5b48cfdbcdb-hpmfb 1/1     Terminating 0          27m   172.18.0.7    minikube   <none>            <none>
kubernetes-bootcamp-5b48cfdbcdb-vslzm 1/1     Terminating 0          27m   172.18.0.8    minikube   <none>            <none>
```

And now we run the get pods to see final state of deployment – the 2 pods that were terminating are truly destroyed.

```
$ kubectl get pods -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP            NODE       NOMINATED NODE   READINESS GATES
kubernetes-bootcamp-5b48cfdbcdb-79n7x 1/1     Running   0          28m   172.18.0.6    minikube   <none>            <none>
kubernetes-bootcamp-5b48cfdbcdb-dhjm8 1/1     Running   0          37m   172.18.0.4    minikube   <none>            <none>
```

Interactive Tutorial - Scaling Your App



Module 6 – Performing a Rolling Update

Objectives: Perform a rolling update using kubectl.

<https://kubernetes.io/docs/tutorials/kubernetes-basics/scale/scale-intro/>

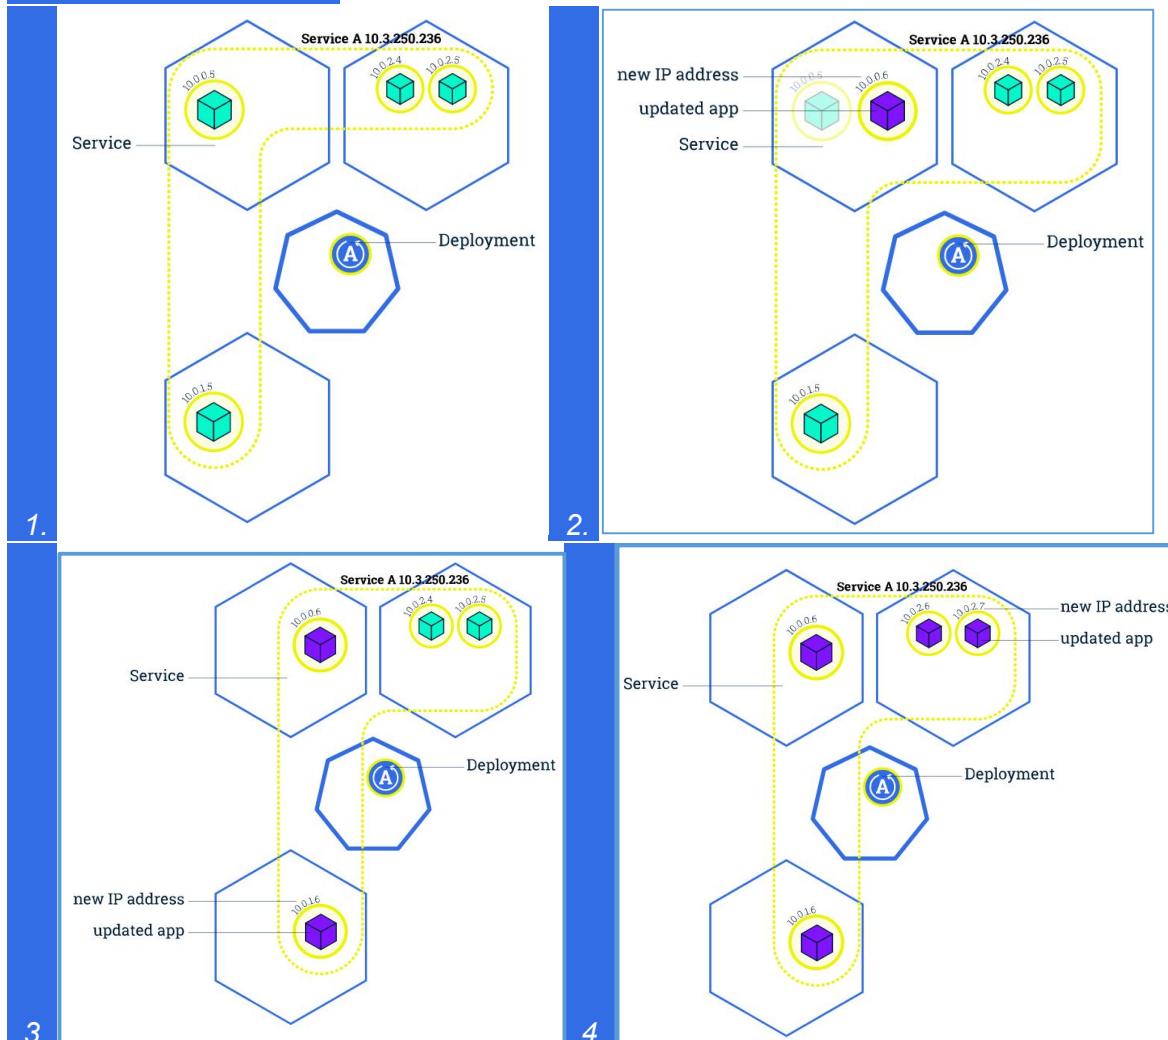
Updating an application

Users expect applications to be available all the time and developers are expected to deploy new versions of them several times a day. In Kubernetes this is done with rolling updates. **Rolling updates** allow Deployments' update to take place with zero downtime by incrementally updating Pods instances with new ones. The new Pods will be scheduled on Nodes with available resources.

In the previous module we scaled our application to run multiple instances. **This is a requirement for performing updates without affecting application availability.** By default, the maximum number of Pods that can be unavailable during the update and the maximum number of new Pods that can be created, is one. Both options can be configured to either numbers or percentages (of Pods). **In Kubernetes, updates are versioned and any Deployment update can be reverted to previous (stable) version.**

Rolling updates overview

Rolling updates allow Deployments' update to take place with zero downtime by incrementally updating Pods instances with new ones.



Similar to application Scaling, if a Deployment is exposed publicly, the Service will load-balance the traffic only to available Pods during the update. An available Pod is an instance that is available to the users of the application.

Rolling updates allow the following actions:

- Promote an application from one environment to another (via container image updates)
- Rollback to previous versions
- Continuous Integration and Continuous Delivery of applications with zero downtime

If a Deployment is exposed publicly, the Service will load-balance the traffic only to available Pods during the update.

In the following interactive tutorial, we'll update our application to a new version, and also perform a rollback.

Lab 6 “Updating Your App” as companion to the Module 6 discussion

The goal of this scenario is to update a deployed application with kubectl set image and to rollback with the rollout undo command.

The online terminal is a pre-configured Linux environment that can be used as a regular console (you can type commands). Clicking on the blocks of code followed by the ENTER sign will execute that command in the terminal.

Kubernetes Bootcamp Terminal

```
$  
$ sleep 1; launch.sh  
Starting Kubernetes. This is expected to take less than a minute  
Kubernetes Started  
$
```

Before any updated, you want to be sure you know your deployments, the running Pods and can confirm the current image version of the application running .

Now, run the kubectl commands to get this information before you update your app:

```
$ kubectl get deployments
```

```
$ kubectl get pods
```

```
$ kubectl get deployments  
NAME                READY   UP-TO-DATE   AVAILABLE   AGE  
kubernetes-bootcamp 4/4     4            4           4m22s  
$ kubectl get pods  
NAME                                READY   STATUS    RESTARTS   AGE  
kubernetes-bootcamp-5b48cfdbc-dtnq5x 1/1     Running   0          4m15s  
kubernetes-bootcamp-5b48cfdbc-xb5p2  1/1     Running   0          4m15s  
kubernetes-bootcamp-5b48cfdbc-xtqbb  1/1     Running   0          4m15s  
kubernetes-bootcamp-5b48cfdbc-z4wv5  1/1     Running   0          4m15s
```

```
$ kubectl describe pods
```

```
$ kubectl describe pods  
Name:                kubernetes-bootcamp-5b48cfdbc-dtnq5x  
Namespace:           default  
Priority:             0  
Node:                minikube/172.17.0.7  
Start Time:          Fri, 17 Jan 2020 05:48:59 +0000  
Labels:               pod-template-hash=5b48cfdbc
```

```
IP: 172.18.0.2
Controlled By: ReplicaSet/kubernetes-bootcamp-5b48cfdcdb
Containers:
  kubernetes-bootcamp:
    Container ID: docker://14f310d3c25505395bef10d91cb48e5dcb5a15973101e83c3fa26e0b65c95cb8
    Image: gcr.io/google-samples/kubernetes-bootcamp:v1
    Image ID: docker-pullable://jocatalin/kubernetes-bootcamp@sha256:0d6b8ee63bb57c5f5b6156f446b3bc3b3c143d233037f3a2f0
0e279c8fcc64af
    Port: 8080/TCP
    Host Port: 0/TCP
    State: Running
      Started: Fri, 17 Jan 2020 05:49:02 +0000
    Ready: True
    Restart Count: 0
    Environment: <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-l89lh (ro)
Conditions:
  Type          Status
  Initialized    True
  Ready          True
  ContainersReady True

  Type: Secret (a volume populated by a Secret)
  SecretName: default-token-l89lh
  Optional: false
QoS Class: BestEffort
Node-Selectors: <none>
Tolerations: node.kubernetes.io/not-ready:NoExecute for 300s
              node.kubernetes.io/unreachable:NoExecute for 300s
Events:
  Type    Reason      Age   From          Message
  ----    -
Normal   Scheduled   4m14s default-scheduler Successfully assigned default/kubernetes-bootcamp-5b48cfdcdb-tnq5x to minikube
Normal   Pulled      4m12s kubelet, minikube Container image "gcr.io/google-samples/kubernetes-bootcamp:v1" already present on machine
Normal   Created     4m12s kubelet, minikube Created container kubernetes-bootcamp
Normal   Started     4m11s kubelet, minikube Started container kubernetes-bootcamp

Name: kubernetes-bootcamp-5b48cfdcdb-xb5p2
Namespace: default
Priority: 0
Node: minikube/172.17.0.7
Start Time: Fri, 17 Jan 2020 05:48:59 +0000

Status: Running
IP: 172.18.0.5
Controlled By: ReplicaSet/kubernetes-bootcamp-5b48cfdcdb
Containers:
  kubernetes-bootcamp:
    Container ID: docker://983efe4995ea7f60d4c212b711424b9dbe77c0bb92de4eaeef16fc0543d8f6ee8
    Image: gcr.io/google-samples/kubernetes-bootcamp:v1
    Image ID: docker-pullable://jocatalin/kubernetes-bootcamp@sha256:0d6b8ee63bb57c5f5b6156f446b3bc3b3c143d233037f3a2f0
0e279c8fcc64af
    Port: 8080/TCP
    Host Port: 0/TCP
    State: Running
      Started: Fri, 17 Jan 2020 05:49:02 +0000
    Ready: True
    Restart Count: 0
    Environment: <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-l89lh (ro)
Conditions:
  Type          Status
  Initialized    True
  Ready          True
```

```
default-token-189lh:
  Type:          Secret (a volume populated by a Secret)
  SecretName:    default-token-189lh
  Optional:      false
QoS Class:       BestEffort
Node-Selectors:  <none>
Tolerations:     node.kubernetes.io/not-ready:NoExecute for 300s
                  node.kubernetes.io/unreachable:NoExecute for 300s
Events:
  Type     Reason      Age   From          Message
  ----     -
Normal    Scheduled   4m14s default-scheduler Successfully assigned default/kubernetes-bootcamp-5b48cfdcdb-xb5p2 to minikube
Normal    Pulled      4m11s kubelet, minikube Container image "gcr.io/google-samples/kubernetes-bootcamp:v1" already present
on machine
Normal    Created     4m11s kubelet, minikube Created container kubernetes-bootcamp
Normal    Started     4m11s kubelet, minikube Started container kubernetes-bootcamp

Name:          kubernetes-bootcamp-5b48cfdcdb-xtqbb
Namespace:     default
Priority:       0
Node:          minikube/172.17.0.7

Annotations:    <none>
Status:         Running
IP:             172.18.0.6
Controlled By:  ReplicaSet/kubernetes-bootcamp-5b48cfdcdb
Containers:
  kubernetes-bootcamp:
    Container ID:  docker://dff7d0dd55724475957383f08f293da71933df92fd6f89877c55285273e8a7db
    Image:         gcr.io/google-samples/kubernetes-bootcamp:v1
    Image ID:      docker-pullable://jocatalin/kubernetes-bootcamp@sha256:0d6b8ee63bb57c5f5b6156f446b3bc3b3c143d233037f3a2f0
0e279c8fcc64af
    Port:         8080/TCP
    Host Port:    0/TCP
    State:        Running
      Started:    Fri, 17 Jan 2020 05:49:02 +0000
    Ready:        True
    Restart Count: 0
    Environment:  <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-189lh (ro)
Conditions:
  Type          Status
  ----          -
Initialized    True

Volumes:
default-token-189lh:
  Type:          Secret (a volume populated by a Secret)
  SecretName:    default-token-189lh
  Optional:      false
QoS Class:       BestEffort
Node-Selectors:  <none>
Tolerations:     node.kubernetes.io/not-ready:NoExecute for 300s
                  node.kubernetes.io/unreachable:NoExecute for 300s
Events:
  Type     Reason      Age   From          Message
  ----     -
Normal    Scheduled   4m14s default-scheduler Successfully assigned default/kubernetes-bootcamp-5b48cfdcdb-xtqbb to minikube
Normal    Pulled      4m12s kubelet, minikube Container image "gcr.io/google-samples/kubernetes-bootcamp:v1" already present
on machine
Normal    Created     4m12s kubelet, minikube Created container kubernetes-bootcamp
Normal    Started     4m11s kubelet, minikube Started container kubernetes-bootcamp

Name:          kubernetes-bootcamp-5b48cfdcdb-z4wv5
Namespace:     default
Priority:       0
```



```
run=kubernetes-bootcamp
Annotations:  <none>
Status:      Running
IP:          172.18.0.4
Controlled By: ReplicaSet/kubernetes-bootcamp-5b48cfdcdb
Containers:
  kubernetes-bootcamp:
    Container ID:  docker://828de037ce548641a8be6f5cc56351011afa2f38d5b6a06f6d22c9beede9e92b
    Image:         gcr.io/google-samples/kubernetes-bootcamp:v1
    Image ID:      docker-pullable://jocatalin/kubernetes-bootcamp@sha256:0d6b8ee63bb57c5f5b6156f446b3bc3b3c143d233037f3a2f0
0e279c8fcc64af
    Port:         8080/TCP
    Host Port:    0/TCP
    State:        Running
      Started:    Fri, 17 Jan 2020 05:49:02 +0000
    Ready:        True
    Restart Count: 0
    Environment:  <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-l89lh (ro)
Conditions:
  Type            Status
  ----            -
  Initialized      True
  Ready            True
  ContainersReady  True
  PodScheduled     True
Volumes:
  default-token-l89lh:
    Type:          Secret (a volume populated by a Secret)
    SecretName:    default-token-l89lh
    Optional:      false
QoS Class:        BestEffort
Node-Selectors:   <none>
Tolerations:      node.kubernetes.io/not-ready:NoExecute for 300s
                  node.kubernetes.io/unreachable:NoExecute for 300s
Events:
  Type    Reason      Age   From          Message
  ----    -
  Normal  Scheduled   4m14s  default-scheduler  Successfully assigned default/kubernetes-bootcamp-5b48cfdcdb-z4wv5 to minikube
  Normal  Pulled      4m12s  kubelet, minikube  Container image "gcr.io/google-samples/kubernetes-bootcamp:v1" already present
on machine
  Normal  Created     4m11s  kubelet, minikube  Created container kubernetes-bootcamp
  Normal  Started     4m11s  kubelet, minikube  Started container kubernetes-bootcamp
$
```

Lab for Step 6.1: Update Your App

Next, to update your app, you will need to update the application to version 2, which you can do with the **set image** command followed by the deployment name and the new image version you wish to deploy.

```
$ kubectl get deployment
```

```
$ kubectl set image deployments/kubernetes-bootcamp kubernetes-bootcamp=jocatalin/kubernetes-bootcamp:v2
```

```
$ kubectl get deployment
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
kubernetes-bootcamp	4/4	4	4	12m

```
$ kubectl set image deployments/kubernetes-bootcamp kubernetes-bootcamp=jocatalin/kubernetes-bootcamp:v2  
deployment.extensions/kubernetes-bootcamp image updated
```

```
$ kubectl get pods
```

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
kubernetes-bootcamp-5b48cfdcdbd-tnq5x	1/1	Terminating	0	12m
kubernetes-bootcamp-5b48cfdcdbd-xb5p2	1/1	Terminating	0	12m
kubernetes-bootcamp-5b48cfdcdbd-xtqbb	1/1	Terminating	0	12m
kubernetes-bootcamp-5b48cfdcdbd-z4wv5	1/1	Terminating	0	12m
kubernetes-bootcamp-cfc74666-2hkl2	1/1	Running	0	10s
kubernetes-bootcamp-cfc74666-lrw97	1/1	Running	0	9s
kubernetes-bootcamp-cfc74666-tjxzc	1/1	Running	0	10s
kubernetes-bootcamp-cfc74666-w6vhl	1/1	Running	0	8s

```
$
```

Wait, then check again

```
$ kubectl get deployment
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
kubernetes-bootcamp	4/4	4	4	13m

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
kubernetes-bootcamp-cfc74666-2hkl2	1/1	Running	0	76s
kubernetes-bootcamp-cfc74666-lrw97	1/1	Running	0	75s
kubernetes-bootcamp-cfc74666-tjxzc	1/1	Running	0	76s
kubernetes-bootcamp-cfc74666-w6vhl	1/1	Running	0	74s

```
$
```

```
$ kubectl get pods -o wide
```

```
$ kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED	NODE	READINESS
GATES									
kubernetes-bootcamp-cfc74666-2hkl2	1/1	Running	0	113s	172.18.0.9	minikube	<none>		<none>
kubernetes-bootcamp-cfc74666-lrw97	1/1	Running	0	112s	172.18.0.11	minikube	<none>		<none>
kubernetes-bootcamp-cfc74666-tjxzc	1/1	Running	0	113s	172.18.0.10	minikube	<none>		<none>
kubernetes-bootcamp-cfc74666-w6vhl	1/1	Running	0	111s	172.18.0.12	minikube	<none>		<none>

```
$
```

Lab for Step 6.2: Verify an Update of Your App

So, now let's check that the app is running and find out the exported IP and exposed Port, we'll start with the **describe service** command once again:

```
$ kubectl get deployments
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
kubernetes-bootcamp 4/4     4            4           17m
$ kubectl describe services/kubernetes-bootcamp | more
Name:                kubernetes-bootcamp
Namespace:           default
Labels:              run=kubernetes-bootcamp
Annotations:         <none>
Selector:            run=kubernetes-bootcamp
Type:               NodePort
IP:                 10.106.128.13
Port:               <unset> 8080/TCP
TargetPort:         8080/TCP
NodePort:           <unset> 30031/TCP
Endpoints:          172.18.0.10:8080,172.18.0.11:8080,172.18.0.12:8080 + 1 more...
Session Affinity:    None
External Traffic Policy: Cluster
Events:             <none>
$
```

Create an environment variable called `NODE_PORT` that has the value of the Node port assigned:

```
export NODE_PORT=$(kubectl get services/kubernetes-bootcamp -o go-template='{{(index .spec.ports 0).nodePort}}') echo NODE_PORT=$NODE_PORT
```

Next, we'll do a `curl` to the the exposed IP and port:

```
curl $(minikube ip):$NODE_PORT
```

We hit a different Pod with every request and we see that all Pods are running the latest version (v2).

```
$ export NODE_PORT=$(kubectl get services/kubernetes-bootcamp -o go-template='{{(index .spec.ports 0).nodePort}}')
$ echo NODE_PORT=$NODE_PORT
NODE_PORT=30031
$ curl $(minikube ip):$NODE_PORT
Hello Kubernetes bootcamp! | Running on: kubernetes-bootcamp-cfc74666-w6vhl | v=2
$ curl $(minikube ip):$NODE_PORT
Hello Kubernetes bootcamp! | Running on: kubernetes-bootcamp-cfc74666-2hk12 | v=2
$ curl $(minikube ip):$NODE_PORT
Hello Kubernetes bootcamp! | Running on: kubernetes-bootcamp-cfc74666-tjxzc | v=2
$ curl $(minikube ip):$NODE_PORT
Hello Kubernetes bootcamp! | Running on: kubernetes-bootcamp-cfc74666-lrw97 | v=2
$
```

Lab for Step 6.2: Rollback an Update of Your App

This time, let's perform another update, and deploy an image tagged as v10. As before, we'll use the **get deployments** command to confirm the state of the deployment before and update the update.

```
$ kubectl get deployments
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
kubernetes-bootcamp 4/4     4            4           22m
$ kubectl set image deployments/kubernetes-bootcamp kubernetes-bootcamp=gcr.io/google-samples/kubernetes-bootcamp:v10
deployment.extensions/kubernetes-bootcamp image updated
$ kubectl get deployments
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
kubernetes-bootcamp 3/4     2            3           22m
$ kubectl get deployments
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
kubernetes-bootcamp 3/4     2            3           22m
```

Something went wrong. Let's check the pods with both **get pods** and **describe pods** to get a better look at what's going on.

```
$ kubectl get deployments
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
kubernetes-bootcamp 3/4     2            3           24m
$ kubectl get pods
NAME                READY   STATUS              RESTARTS   AGE
kubernetes-bootcamp-547469f5dd-6rf59 0/1     ErrImagePull        0          100s
kubernetes-bootcamp-547469f5dd-8x26j 0/1     ImagePullBackOff    0          100s
kubernetes-bootcamp-cfc74666-2hkl2    1/1     Running             0          11m
kubernetes-bootcamp-cfc74666-lrw97    1/1     Running             0          11m
kubernetes-bootcamp-cfc74666-tjxzc    1/1     Running             0          11m
$
```

\$ kubectl describe pods | more

```
$ kubectl describe pods | more
Name:                kubernetes-bootcamp-547469f5dd-6rf59
Namespace:           default
Priority:             0
Node:                minikube/172.17.0.7
Start Time:          Fri, 17 Jan 2020 06:11:36 +0000
Labels:              pod-template-hash=547469f5dd
                    run=kubernetes-bootcamp
Annotations:         <none>
Status:              Pending
IP:                  172.18.0.4
Controlled By:        ReplicaSet/kubernetes-bootcamp-547469f5dd
Containers:
  kubernetes-bootcamp:
    Container ID:      gcr.io/google-samples/kubernetes-bootcamp:v10
    Image ID:          
    Port:              8080/TCP
    Host Port:         0/TCP
    State:              Waiting
    Reason:             ImagePullBackOff
    Ready:             False
  Restart Count:      0
  Environment:        <none>
  Mounts:
    /var/run/secrets/kubernetes.io/serviceaccount from default-token-l89lh (ro)
Conditions:
  Type              Status
  Initialized        True
  Ready              False
  ContainersReady    False
  PodScheduled       True
Volumes:
  default-token-l89lh:
    Type:             Secret (a volume populated by a Secret)
    SecretName:        default-token-l89lh
    Optional:          false
QoS Class:           BestEffort
Node-Selectors:      <none>
Tolerations:         node.kubernetes.io/not-ready:NoExecute for 300s
                    node.kubernetes.io/unreachable:NoExecute for 300s
Events:
  Type              Reason              Age              From              Message
  ----              -
```

```
Normal    Scheduled    2m41s    default-scheduler    Successfully assigned default/kubernetes-bootcamp-547469f5dd-6r
f59 to minikube
Normal    Pulling      81s (x4 over 2m40s)    kubelet, minikube    Pulling image "gcr.io/google-samples/kubernetes-bootcamp:v10"
Warning   Failed       81s (x4 over 2m39s)    kubelet, minikube    Failed to pull image "gcr.io/google-samples/kubernetes-bootcamp
:v10": rpc error: code = Unknown desc = Error response from daemon: manifest for gcr.io/google-samples/kubernetes-bootcamp:v1
0 not found
Warning   Failed       81s (x4 over 2m39s)    kubelet, minikube    Error: ErrImagePull
Normal    BackOff      53s (x6 over 2m39s)    kubelet, minikube    Back-off pulling image "gcr.io/google-samples/kubernetes-bootca
mp:v10"
Warning   Failed       42s (x7 over 2m39s)    kubelet, minikube    Error: ImagePullBackOff
```

Above – the events tell us a lot – and we know why we have 2 pods (original and new) in a bad state

```
Name:      kubernetes-bootcamp-547469f5dd-8x26j
Namespace: default
Priority:   0
Node:      minikube/172.17.0.7
Start Time: Fri, 17 Jan 2020 06:11:36 +0000
Labels:    pod-template-hash=547469f5dd
           run=kubernetes-bootcamp
Annotations: <none>

Status:    Pending
IP:        172.18.0.2
Controlled By: ReplicaSet/kubernetes-bootcamp-547469f5dd
Containers:
  kubernetes-bootcamp:
    Container ID:
    Image:      gcr.io/google-samples/kubernetes-bootcamp:v10
    Image ID:
    Port:       8080/TCP
    Host Port:  0/TCP
    State:      Waiting
      Reason:   ImagePullBackOff
    Ready:      False
    Restart Count: 0
    Environment: <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-l89lh (ro)
Conditions:
  Type           Status
  Initialized     True
  Ready          False
  ContainersReady False
  PodScheduled    True
Volumes:
  default-token-l89lh:
    Type:      Secret (a volume populated by a Secret)
    SecretName: default-token-l89lh
    Optional:   false
QoS Class:     BestEffort
Node-Selectors: <none>
Tolerations:   node.kubernetes.io/not-ready:NoExecute for 300s
               node.kubernetes.io/unreachable:NoExecute for 300s
```

Events:

Type	Reason	Age	From	Message
Normal	Scheduled	2m41s	default-scheduler	Successfully assigned default/kubernetes-bootcamp-547469f5dd-8x26j to minikube
Normal	Pulling	69s (x4 over 2m40s)	kubelet, minikube	Pulling image "gcr.io/google-samples/kubernetes-bootcamp:v10"
Warning	Failed	69s (x4 over 2m40s)	kubelet, minikube	Failed to pull image "gcr.io/google-samples/kubernetes-bootcamp:v10": rpc error: code = Unknown desc = Error response from daemon: manifest for gcr.io/google-samples/kubernetes-bootcamp:v10 not found
Warning	Failed	69s (x4 over 2m40s)	kubelet, minikube	Error: ErrImagePull
Normal	BackOff	56s (x6 over 2m39s)	kubelet, minikube	Back-off pulling image "gcr.io/google-samples/kubernetes-bootcamp:v10"
Warning	Failed	41s (x7 over 2m39s)	kubelet, minikube	Error: ImagePullBackOff

Name: kubernetes-bootcamp-cfc74666-2hk12
Namespace: default
Priority: 0
Node: minikube/172.17.0.7
Start Time: Fri, 17 Jan 2020 06:01:24 +0000
Labels: pod-template-hash=cfc74666
run=kubernetes-bootcamp
Annotations: <none>
Status: Running
IP: 172.18.0.9
Controlled By: ReplicaSet/kubernetes-bootcamp-cfc74666
Containers:
 kubernetes-bootcamp:
 Container ID: docker://59dff6b522615e446e3a5d31107de1302da523653a2ad2fad1c5f226edce28df
 Image: jocatalin/kubernetes-bootcamp:v2
 Image ID: docker-pullable://jocatalin/kubernetes-bootcamp@sha256:fb1a3ced00cecf1f83f18ab5cd14199e30adc1b49aa4244f5

d65ad3f5feb2a5

Port: 8080/TCP
Host Port: 0/TCP
State: Running
 Started: Fri, 17 Jan 2020 06:01:25 +0000
Ready: True
Restart Count: 0
Environment: <none>
Mounts:
 /var/run/secrets/kubernetes.io/serviceaccount from default-token-l89lh (ro)

Conditions:

Type	Status
Initialized	True
Ready	True
ContainersReady	True
PodScheduled	True

Volumes:

default-token-l89lh:
 Type: Secret (a volume populated by a Secret)
 SecretName: default-token-l89lh
 Optional: false

.....

```
Type:                Status
Initialized          True
Ready                True
ContainersReady      True
PodScheduled         True

Volumes:
default-token-l89lh:
  Type:              Secret (a volume populated by a Secret)
  SecretName:        default-token-l89lh
  Optional:          false
QoS Class:           BestEffort
Node-Selectors:      <none>
Tolerations:         node.kubernetes.io/not-ready:NoExecute for 300s
                    node.kubernetes.io/unreachable:NoExecute for 300s

Events:
  Type    Reason      Age   From          Message
  ----    -
Normal    Scheduled   19m   default-scheduler   Successfully assigned default/kubernetes-bootcamp-cfc74666-tjxzc to minikube
Normal    Pulled      19m   kubelet, minikube   Container image "jocatalin/kubernetes-bootcamp:v2" already present on machine
Normal    Created     19m   kubelet, minikube   Created container kubernetes-bootcamp
Normal    Started     19m   kubelet, minikube   Started container kubernetes-bootcamp

$
```

Use “rollback undo” command

```
$ kubectl get pods
NAME                                READY   STATUS              RESTARTS   AGE
kubernetes-bootcamp-547469f5dd-6rf59 0/1     ImagePullBackOff    0          11m
kubernetes-bootcamp-547469f5dd-8x26j 0/1     ErrImagePull        0          11m
kubernetes-bootcamp-cfc74666-2hkl2    1/1     Running             0          21m
kubernetes-bootcamp-cfc74666-lrw97    1/1     Running             0          21m
kubernetes-bootcamp-cfc74666-tjxzc    1/1     Running             0          21m

$ kubectl rollout undo deployments/kubernetes-bootcamp
deployment.extensions/kubernetes-bootcamp rolled back

$ kubectl get pods
NAME                                READY   STATUS              RESTARTS   AGE
kubernetes-bootcamp-547469f5dd-6rf59 0/1     Terminating       0          11m
kubernetes-bootcamp-547469f5dd-8x26j 0/1     Terminating       0          11m
kubernetes-bootcamp-cfc74666-2hkl2    1/1     Running             0          21m
kubernetes-bootcamp-cfc74666-lrw97    1/1     Running             0          21m
kubernetes-bootcamp-cfc74666-sfjpm 0/1     ContainerCreating   0          3s
kubernetes-bootcamp-cfc74666-tjxzc    1/1     Running             0          21m

$
```

Wait and check again to give the Terminating Pods time to clear, and the new Container to be created

```
$ kubectl get pods
NAME                                READY   STATUS              RESTARTS   AGE
kubernetes-bootcamp-cfc74666-2hkl2    1/1     Running             0          22m
kubernetes-bootcamp-cfc74666-lrw97    1/1     Running             0          22m
kubernetes-bootcamp-cfc74666-sfjpm 1/1     Running             0          47s
kubernetes-bootcamp-cfc74666-tjxzc    1/1     Running             0          22m

$
```

<https://kubernetes.io/docs/tutorials/kubernetes-basics/>

Four Pods are running. Check again the image deployed on them:

```
kubectl describe pods
```

```
...
$ kubectl describe pods | grep v2
Image:          jocatalin/kubernetes-bootcamp:v2
Normal Pulled   25m kubelet, minikube Container image "jocatalin/kubernetes-bootcamp:v2" already present on machine
Image:          jocatalin/kubernetes-bootcamp:v2
Normal Pulled   25m kubelet, minikube Container image "jocatalin/kubernetes-bootcamp:v2" already present on machine
Image:          jocatalin/kubernetes-bootcamp:v2
Normal Pulled   3m36s kubelet, minikube Container image "jocatalin/kubernetes-bootcamp:v2" already present on machine
Image:          jocatalin/kubernetes-bootcamp:v2
Normal Pulled   25m kubelet, minikube Container image "jocatalin/kubernetes-bootcamp:v2" already present on machine
$
```

We see that the deployment is using a stable version of the app (v2). The Rollback was successful.

SUCCESS – ALL OBJECTIVES MET

Tutorials

Hello Minikube

▼ Learn Kubernetes Basics

- Learn Kubernetes Basics
- ▶ Create a Cluster
- ▶ Deploy an App
- ▶ Explore Your App
- ▶ Expose Your App Publicly
- ▶ Scale Your App
- ▶ Update Your App



Tutorials

Hello Minikube

▼ Learn Kubernetes Basics

- Learn Kubernetes Basics
- ▶ Create a Cluster
- ▶ Deploy an App
- ▶ Explore Your App
- ▶ Expose Your App Publicly
- ▶ Scale Your App
- ▶ Update Your App

APPENDIX – Notes during Module 2

What part of the API Server is exposed?

NOTE: Finding out the allowed “paths” of the API is easy – type anything that will not be normal to be there (like ‘curl’)

```
$ curl http://localhost:8001/curl
{
  "paths": [
    "/apis",
    "/apis/",
    "/apis/apiextensions.k8s.io",
    "/apis/apiextensions.k8s.io/v1beta1",
    "/healthz",
    "/healthz/etcd",
    "/healthz/log",
    "/healthz/ping",
    "/healthz/poststarthook/crd-informer-synced",
    "/healthz/poststarthook/generic-apiserver-start-informers",
    "/healthz/poststarthook/start-apiextensions-controllers",
    "/healthz/poststarthook/start-apiextensions-informers",
    "/metrics",
    "/openapi/v2",
    "/version"
  ]
}
```

You can see why /version provided the earlier output.

```
$ curl http://localhost:8001/healthz/etcd
ok$
$
```

Some output, like /metrics will provide more output than can be captured here, but you can pipe that output in to grep (e.g., “ curl <http://localhost:8001/metrics> | grep etcd | wc -l “ and you can see there is a lot of output still...

```
$ curl http://localhost:8001/metrics | grep etcd | wc -l
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total     Spent    Left  Speed
100 2557k    0 2557k    0    0  26.2M      0  --:--:-- --:--:-- --:--:-- 26.2M
2827
$
```

<https://kubernetes.io/docs/tutorials/kubernetes-basics/>

Appendix – Notes on Module 3

Enter container with Bash and exit

```
$ kubectl exec -ti $POD_NAME bash
root@kubernetes-bootcamp-5b48cfdcdb-vnnzx:/# ls
bin boot core dev etc home lib lib64 media mnt opt proc root run sbin server.js srv sys tmp usr var
root@kubernetes-bootcamp-5b48cfdcdb-vnnzx:/# exit
exit
$
```

```
$ kubectl logs $POD_NAME
Kubernetes Bootcamp App Started At: 2020-01-14T03:32:34.740Z | Running On: kubernetes-bootcamp-5b48cfdcdb-vnnzx

Running On: kubernetes-bootcamp-5b48cfdcdb-vnnzx | Total Requests: 1 | App Uptime: 1361.899 seconds | Log Time: 2020-01-14T03:55:16.639Z
Running On: kubernetes-bootcamp-5b48cfdcdb-vnnzx | Total Requests: 2 | App Uptime: 2234.447 seconds | Log Time: 2020-01-14T04:09:49.187Z
$ kubectl exec -ti $POD_NAME bash
root@kubernetes-bootcamp-5b48cfdcdb-vnnzx:/# ls
bin boot core dev etc home lib lib64 media mnt opt proc root run sbin server.js srv sys tmp usr var
root@kubernetes-bootcamp-5b48cfdcdb-vnnzx:/# exit
exit
$
```

```
root@kubernetes-bootcamp-5b48cfdcdb-vnnzx:/# cat server.js
var http = require('http');
var requests=0;
var podname= process.env.HOSTNAME;
var startTime;
var host;
var handleRequest = function(request, response) {
  response.setHeader('Content-Type', 'text/plain');
  response.writeHead(200);
  response.write("Hello Kubernetes bootcamp! | Running on: ");
  response.write(host);
  response.end(" | v=1\n");
  console.log("Running On:" ,host, "| Total Requests:", ++requests,"| App Uptime:", (new Date() - startTime)/1000 , "seconds", "| Log Time:",new Date());
}
var www = http.createServer(handleRequest);
www.listen(8080,function () {
  startTime = new Date();
  host = process.env.HOSTNAME;
  console.log ("Kubernetes Bootcamp App Started At:",startTime, "| Running On: " ,host, "\n" );
});
root@kubernetes-bootcamp-5b48cfdcdb-vnnzx:/#
```

```
root@kubernetes-bootcamp-5b48cfdcdb-vnnzx:/# curl localhost:8080
Hello Kubernetes bootcamp! | Running on: kubernetes-bootcamp-5b48cfdcdb-vnnzx | v=1
root@kubernetes-bootcamp-5b48cfdcdb-vnnzx:/#
```